

Intro to Machine Learning Final Project

Enron Dataset

This dataset consists of 146 employees of Enron and 21 features consisting of both financial and email information. 18 of these 146 employees have been identified as “Person of Interest” or a “POI”. Two of these “employees” were determined to be outliers and were eliminated for this analysis. One was the key “TOTAL” which is an artifact of an earlier dataset where all the columns were being summed. The other is “THE TRAVEL AGENCY IN THE PARK” as it was not a valid employee. I investigated how many missing values there were in the dataset and came back with the follow results:

```
{'bonus': 63,
 'deferral_payments': 106,
 'deferred_income': 96,
 'director_fees': 128,
 'email_address': 33,
 'exercised_stock_options': 43,
 'expenses': 50,
 'from_messages': 58,
 'from_poi_to_this_person': 58,
 'from_this_person_to_poi': 58,
 'loan_advances': 141,
 'long_term_incentive': 79,
 'other': 53,
 'restricted_stock': 35,
 'restricted_stock_deferred': 127,
 'salary': 50,
 'shared_receipt_with_poi': 58,
 'to_messages': 58,
 'total_payments': 21,
 'total_stock_value': 19}
```

The features with the most amount of ‘NaN’ values were director fees, loan advances, and restricted stock deferred. This make sense as these values probably only apply to higher ranking employees.

Algorithm

Each of the models was run through a testing function to properly test the new algorithm parameters. This function uses a k-folds cross validation program to split the data into training and testing sets many times to test the functionality of the model. For most of these tests, 1000 different sets of the data were created to validate. Validation over many sets of data is important because if we were to test the model once over one training set, we may get inaccurate results, especially with the limited number of POIs. We want to test the model many, many times and take a look at the results to ensure they are accurate. This function also included an optional variable that would implement a function to scale the features of the dataset. This function takes each feature and scales it down to a value of 0 to 1. This ensures that all features are treated equally. This is important in datasets like the Enron dataset, because features like finances and quantity of email have such different scales. I ran model through with and without scaled features to test impact.

Model	Score	Precision	Recall
Decision Tree Classifier (no scaling)	.81273	.2395	.2322
Decision Tree Classifier (scaled features)	.81306	.2222	.2290
Gaussian Naïve Bayes (no scaling)	.33433	.1474	.8345
Gaussian Naïve Bayes (scaled features)	.56533	.2192	.7250
SVC (no scaling)	.86666	0.000	0.000
SVC (scaled features)	.86586	.0185	.0095

While the Naïve Bayes model has a significantly higher recall, it was identifying far too many false positives. I decided to use the Decision Tree model as it has a far more balanced score. The feature scaling had a noticeable impact on both the Gaussian Naïve Bayes and SVC models, but not the Decision Tree.

Features

I decided to add two new features to the dataset, each a fraction of emails sent to or from that employee and a person of interest. I thought this would be a more interesting feature than just seeing how many emails were sent to POI's and from POI's, we would instead be able to look at what fraction of their total emails were made up with communication with POI's.

These features gave a slight boost to the precision and recall of the Decision Tree model.

	Precision	Recall
Before new features	.2395	.2322
With new features	.2459	.2475

The fraction of emails sent to a POI proved to be an important feature in the final feature selection.

Because I decided to use the Decision Tree model, I also used a built in function of that model to select features (feature_importances_). This function calculates the importance of each feature as the normalized reduction of the weight of that feature on the model. I used a function that would create the Decision Tree model, weigh the features, and then trim any features that had scored a zero (not important). This function was run until the amount of features was reduced to only features that had importance greater than zero. The only features that remained were: 'total_payments', 'shared_receipt_with_poi', 'fraction_to_poi', 'exercised_stock_options', 'expenses', 'restricted_stock'.

	Precision	Recall
All Features	.2459	.2475
Reduced Features	.3620	.3800

This had an extremely large impact on the model with a much higher precision and recall.

Algorithm Tuning

Each algorithm uses a set of parameters that can have drastic results on the final result. It's important to test each algorithm with a variety of parameters to ensure that you are getting the highest results possible. Thankfully, there are ways to automate this process to ensure you can test a large group of

parameters programmatically. I tuned the algorithm using the Grid Search CV function, while weighting the scoring toward recall, since I wanted to model to favor identifying POIs. Using the Grid Search CV function is essential as it allows us to test many attributes to make sure that the model is properly tuned with the correct attributes. The parameter grid used for searching was determined by hand testing variables to get a good feel for an appropriate range. The parameters I used for testing were:

'max_features': ['auto', 'sqrt', 'log2']

Max features determines the amount of features to use when splitting. 'auto' and 'sqrt' will test by taking the max features as the square root of the features list of the data set. 'log2' will take the log2 of the list of features.

'min_samples_split': [2, 3, 4, 5, 6, 7, 8, 9, 10]

Minimum samples split will decide how many samples of the feature need to be present for the tree to split based on this feature.

'max_depth' : [None, 2, 4, 6, 8, 10]

Max depth is the depth of the decision tree. This will determine how many splits the tree will make.

The parameters returned by the search are:

'max_features'	'sqrt'
'min_samples_split'	3
'max_depth'	None

The entire set of parameters returned by the search is:

```
DecisionTreeClassifier(compute_importances=None, criterion='gini',  
                        max_depth=None, max_features='sqrt', max_leaf_nodes=None,  
                        min_density=None, min_samples_leaf=1, min_samples_split=3,  
                        random_state=None, splitter='best')
```

After the algorithm tuning the model performed even better:

	Precision	Recall
Before Tuning	.3620	.3800
After Tuning	.3705	.3873

Validation

What we are looking for more than accuracy is a high precision and recall score. The precision identifies the number of true positives we've identified over the all positives. What this means is how many of the employees as a POI / non-POI correctly. Next we have the recall score is the number of true positives divided by true positives added to the false negatives. This is essentially how many employees were identified as a POI and actually are a POI. Because we would rather get all POIs while possibly incorrectly

flagging some innocent people as POIs, we prefer a higher recall. While the accuracy went slightly down, the precision and recall of the model both rose to .392 and .366. The model is now much better at identifying POIs from the selected features.